

```
In [ ]: import torch
        from torchvision import datasets, transforms, models
        import torchvision.utils as vutils
        from torch.utils.data import DataLoader
        import torch.nn as nn
        import torch.nn.functional as F
        from sklearn.metrics import confusion_matrix
        import matplotlib.pyplot as plt
        import numpy as np
        from torch.utils.data.dataset import random_split
        from torchsummary import summary
        import time
        import seaborn as sns
        # import seaborn as sns

        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        # device = "cpu"
```

```
In [ ]: device
```

```
Out[ ]: device(type='cuda')
```

I have provided visualizations at the bottom of the notebook. I have included some explanations for how I planned & constructed my model

```
In [ ]: transform = transforms.Compose([
        transforms.ToTensor(), # Convert images to PyTorch tensors
        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)), # Normalize the images
        transforms.RandomHorizontalFlip(), # Randomly flip the images horizontally with
        transforms.RandomRotation(10),
        transforms.GaussianBlur()
    ])

    train_data = datasets.CIFAR10(
        root='data',
        train=True,
        download=True,
        transform=transform
    )

    test_data = datasets.CIFAR10(
        root='data',
        train=False,
        download=True,
        transform=transform
    )

    batch_size = 64
    total_size = len(train_data)

    # Determine the split ratio
    train_ratio = 0.9
```

```

validation_ratio = 1 - train_ratio

# Calculate Lengths
train_length = int(total_size * train_ratio)
validation_length = total_size - train_length # ensures no data point is left out

# Split the dataset
train_dataset, valid_dataset = random_split(train_data, lengths=[train_length, validation_length])
# train_dataset, valid_dataset = random_split(train_data, lengths=[55000, 5000])
trainloader = DataLoader(train_data, batch_size=50)
testloader = DataLoader(test_data, batch_size=50)

valid_loader = DataLoader(
    dataset=valid_dataset,
    batch_size=batch_size,
    drop_last=False,
    num_workers=3,
    shuffle=False
)

```

Files already downloaded and verified

Files already downloaded and verified

I decided to experiment with modularizing my convolutions. I defined a 'longblock' class as a container for my deep convolutions. 'Shortblock' is for shorter but wider convolutions (bigger kernel and less channels)

After experimenting, I found that adding skip connections vastly improved performance. The original model had the long & shortblocks, as well as batch norm. The test performance wasn't good, so I added dropout and it brought it above 80%

```

In [ ]: class Longblock(nn.Module):
        """
        Takes input with 3 channels. Follows shortblock
        Outputs
        """
        def __init__(self, in_channels, out_channels, first=False):
            super(Longblock, self).__init__()
            self.first = first
            #Test convolution, just adds more channels
            self.conv1 = nn.Conv2d(in_channels=in_channels, out_channels=128, kernel_size=3, padding=1)
            self.conv2 = nn.Conv2d(in_channels=128, out_channels=160, kernel_size=3, padding=1)
            self.conv3 = nn.Conv2d(in_channels=160, out_channels=out_channels, kernel_size=3, padding=1)
            self.bn1 = nn.BatchNorm2d(num_features=128)
            self.bn2 = nn.BatchNorm2d(num_features=160)
            # self.bn3 = nn.BatchNorm2d(num_features=3)
            self.adjust_channels = nn.Conv2d(in_channels=in_channels, out_channels=out_channels, kernel_size=1, padding=0)
            self.in_channels = in_channels
            self.out_channels = out_channels
            self.dropout = nn.Dropout(p=0.2)
        def forward(self, x):
            #For every step, apply RELU & Batch norm
            identity = x
            identity = self.adjust_channels(identity)
            x = F.relu(self.bn1(self.conv1(x)))

```

```

x = F.relu(self.bn2(self.conv2(x)))
# x = F.relu(self.bn3(self.conv3(x)))
x = self.conv3(x)
# x = F.max_pool2d(x, kernel_size=2)
x += identity
x = self.dropout(x)

return x

```

```
class Shortblock(nn.Module):
```

```
    """
```

```
    Takes input with 3 channels
    Fatter than the longblock.
    """
```

```

def __init__(self, in_channels, out_channels):
    super(Shortblock, self).__init__()
    self.conv1 = nn.Conv2d(in_channels=in_channels, out_channels=20, kernel_size=3, padding=1)
    self.conv2 = nn.Conv2d(in_channels=20, out_channels=64, kernel_size=7, padding=3)
    self.conv3 = nn.Conv2d(in_channels=64, out_channels=out_channels, kernel_size=3, padding=1)
    self.bn1 = nn.BatchNorm2d(num_features=20)
    self.bn2 = nn.BatchNorm2d(num_features=64)
    self.bn3 = nn.BatchNorm2d(num_features=out_channels)
    self.dropout = nn.Dropout(p=0.2)

```

```

def forward(self, x):
    #For every step, apply RELU & Batch norm. Then dropout and maxpool
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.bn2(self.conv2(x)))
    x = F.relu(self.bn3(self.conv3(x)))
    x = self.dropout(x)
    x = F.max_pool2d(x, kernel_size=3)
    return x

```

```
class Model(nn.Module):
```

```
    def __init__(self):
```

```

        super(Model, self).__init__()
        #Resizes to 3 channels from 1
        # self.resizeconv = nn.Conv2d(in_channels=3, out_channels=3, kernel_size=3, padding=1)
        self.shortblock = Shortblock(3, 128)
        self.longblock1 = Longblock(128, 128)
        self.longblock2 = Longblock(128, 256)
        self.longblock3 = Longblock(256, 512)
        self.longblock4 = Longblock(512, 256)
        self.shortblock2 = Shortblock(256, 128)
        self.longblock5 = Longblock(128, 64)

        self.Linear1 = nn.Linear(576, 100)
        self.Linear2 = nn.Linear(100, 500)
        self.Linear3 = nn.Linear(500, 10)

```

```
    def forward(self, x):
```

```

        #Convolution Blocks
        #This totals 10 convolutions. Resize is to add more channels initially, the
        # x = self.resizeconv(x)
        x = self.shortblock(x)

```

```
x = self.longblock1(x)
x = self.longblock2(x)
x = self.longblock3(x)
x = self.longblock4(x)
x = self.shortblock2(x)
x = self.longblock5(x)
self.dropout = nn.Dropout(0.2)
# print(x.shape)

#Flatten and run through Linear Layers
x = torch.flatten(x, start_dim=1)
x = F.relu(self.Linear1(x))
x = F.relu(self.Linear2(x))
x = self.dropout(x)
x = self.Linear3(x)
return x
```

```
In [ ]: print(Model())
```

```

Model(
  (shortblock): Shortblock(
    (conv1): Conv2d(3, 20, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
    (conv2): Conv2d(20, 64, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
    (conv3): Conv2d(64, 128, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
    (bn1): BatchNorm2d(20, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (dropout): Dropout(p=0.2, inplace=False)
  )
  (longblock1): Longblock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv2): Conv2d(128, 160, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(160, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn2): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (adjust_channels): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
    (dropout): Dropout(p=0.2, inplace=False)
  )
  (longblock2): Longblock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv2): Conv2d(128, 160, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(160, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn2): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (adjust_channels): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1))
    (dropout): Dropout(p=0.2, inplace=False)
  )
  (longblock3): Longblock(
    (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv2): Conv2d(128, 160, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(160, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn2): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (adjust_channels): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1))
    (dropout): Dropout(p=0.2, inplace=False)
  )
  (longblock4): Longblock(
    (conv1): Conv2d(512, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv2): Conv2d(128, 160, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(160, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (bn2): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (adjust_channels): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (dropout): Dropout(p=0.2, inplace=False)
  )
)

```

```

)
(shortblock2): Shortblock(
  (conv1): Conv2d(256, 20, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
  (conv2): Conv2d(20, 64, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
  (conv3): Conv2d(64, 128, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
  (bn1): BatchNorm2d(20, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (dropout): Dropout(p=0.2, inplace=False)
)
(longblock5): Longblock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(128, 160, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(160, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (bn2): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (adjust_channels): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1))
  (dropout): Dropout(p=0.2, inplace=False)
)
  (Linear1): Linear(in_features=576, out_features=100, bias=True)
  (Linear2): Linear(in_features=100, out_features=500, bias=True)
  (Linear3): Linear(in_features=500, out_features=10, bias=True)
)

```

```
In [ ]: summary(Model().to(device),(3,32,32))
```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 20, 32, 32]	2,960
BatchNorm2d-2	[-1, 20, 32, 32]	40
Conv2d-3	[-1, 64, 32, 32]	62,784
BatchNorm2d-4	[-1, 64, 32, 32]	128
Conv2d-5	[-1, 128, 32, 32]	401,536
BatchNorm2d-6	[-1, 128, 32, 32]	256
Dropout-7	[-1, 128, 32, 32]	0
Shortblock-8	[-1, 128, 10, 10]	0
Conv2d-9	[-1, 128, 10, 10]	16,512
Conv2d-10	[-1, 128, 10, 10]	147,584
BatchNorm2d-11	[-1, 128, 10, 10]	256
Conv2d-12	[-1, 160, 10, 10]	184,480
BatchNorm2d-13	[-1, 160, 10, 10]	320
Conv2d-14	[-1, 128, 10, 10]	184,448
Dropout-15	[-1, 128, 10, 10]	0
Longblock-16	[-1, 128, 10, 10]	0
Conv2d-17	[-1, 256, 10, 10]	33,024
Conv2d-18	[-1, 128, 10, 10]	147,584
BatchNorm2d-19	[-1, 128, 10, 10]	256
Conv2d-20	[-1, 160, 10, 10]	184,480
BatchNorm2d-21	[-1, 160, 10, 10]	320
Conv2d-22	[-1, 256, 10, 10]	368,896
Dropout-23	[-1, 256, 10, 10]	0
Longblock-24	[-1, 256, 10, 10]	0
Conv2d-25	[-1, 512, 10, 10]	131,584
Conv2d-26	[-1, 128, 10, 10]	295,040
BatchNorm2d-27	[-1, 128, 10, 10]	256
Conv2d-28	[-1, 160, 10, 10]	184,480
BatchNorm2d-29	[-1, 160, 10, 10]	320
Conv2d-30	[-1, 512, 10, 10]	737,792
Dropout-31	[-1, 512, 10, 10]	0
Longblock-32	[-1, 512, 10, 10]	0
Conv2d-33	[-1, 256, 10, 10]	131,328
Conv2d-34	[-1, 128, 10, 10]	589,952
BatchNorm2d-35	[-1, 128, 10, 10]	256
Conv2d-36	[-1, 160, 10, 10]	184,480
BatchNorm2d-37	[-1, 160, 10, 10]	320
Conv2d-38	[-1, 256, 10, 10]	368,896
Dropout-39	[-1, 256, 10, 10]	0
Longblock-40	[-1, 256, 10, 10]	0
Conv2d-41	[-1, 20, 10, 10]	250,900
BatchNorm2d-42	[-1, 20, 10, 10]	40
Conv2d-43	[-1, 64, 10, 10]	62,784
BatchNorm2d-44	[-1, 64, 10, 10]	128
Conv2d-45	[-1, 128, 10, 10]	401,536
BatchNorm2d-46	[-1, 128, 10, 10]	256
Dropout-47	[-1, 128, 10, 10]	0
Shortblock-48	[-1, 128, 3, 3]	0
Conv2d-49	[-1, 64, 3, 3]	8,256
Conv2d-50	[-1, 128, 3, 3]	147,584
BatchNorm2d-51	[-1, 128, 3, 3]	256
Conv2d-52	[-1, 160, 3, 3]	184,480
BatchNorm2d-53	[-1, 160, 3, 3]	320

Conv2d-54	[-1, 64, 3, 3]	92,224
Dropout-55	[-1, 64, 3, 3]	0
Longblock-56	[-1, 64, 3, 3]	0
Linear-57	[-1, 100]	57,700
Linear-58	[-1, 500]	50,500
Linear-59	[-1, 10]	5,010

=====

Total params: 5,622,542
Trainable params: 5,622,542
Non-trainable params: 0

Input size (MB): 0.01
Forward/backward pass size (MB): 10.18
Params size (MB): 21.45
Estimated Total Size (MB): 31.64

```
In [ ]: random_seed = 1
batch_size = 256
learning_rate = 0.03
num_epochs = 15
num_classes = 10
def accuracy(model, data_loader, device):
    with torch.no_grad():
        model = model.train()
        true_pred = 0
        tot_samples = 0
        for imgs, labels in data_loader:
            imgs = imgs.to(device)
            labels = labels.to(device)
            logits = model(imgs)
            _, label_pred = torch.max(logits, axis=1)
            true_pred += (label_pred==labels).sum()
            tot_samples += labels.shape[0]
        acc = (true_pred/float(tot_samples))*100
    return acc
```

```
In [ ]: torch.manual_seed(random_seed)
model = Model()
model = model.to(device)

optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=2, gamma=0.5)
#For storing data on losses
train_loss= []

start = time.time()
for epoch in range(num_epochs):
    model = model.train()
    for batch_idx, (imgs, labels) in enumerate(trainloader):
        imgs = imgs.to(device)
        labels = labels.to(device)

        ## Forward Propagation - extract features and classify
        logits = model(imgs)
        loss = F.cross_entropy(logits, labels)
```



```

        #zero out the gradients
        optimizer.zero_grad()
        #estimate new gradients
        loss.backward()
        #update parameters
        optimizer.step()

    if not (batch_idx + 1) % 100:
        print(
            f"Epoch: {epoch + 1:03d}/{num_epochs:03d} | "
            f"Batch: {batch_idx + 1:03d}/{len(trainloader):03d} | "
            f"Loss: {loss:.4f}"
        )

# Tracking the Learning Rate Scheduler
prev_lr = optimizer.param_groups[0]["lr"]
scheduler.step()
current_lr = optimizer.param_groups[0]["lr"]
print(f"Epoch: {epoch+1:03d} Learning Rate {prev_lr:.8f} -> {current_lr:.8f}")

# Evaluate Performance after each epoch
model = model.eval()
tr_acc = accuracy(model, trainloader, device)
valid_acc = accuracy(model, valid_loader, device)
print(f"Train Accuracy: {tr_acc:0.3f}")
print(f"Validation Accuracy: {valid_acc:0.3f}")
if epoch % 5 == 0:
    ts_acc = accuracy(model, testloader, device)
    print(f"Test Accuracy: {ts_acc:0.3f}")

print(f"Time elapsed so far: {(time.time() - start) / 60:.2f} min")
train_loss.append(loss)
print(f"Total Train Time: {(time.time() - start) / 60:.2f} min")

```

Epoch: 001/015 | Batch: 100/1000 | Loss: 1.8576
Epoch: 001/015 | Batch: 200/1000 | Loss: 1.7261
Epoch: 001/015 | Batch: 300/1000 | Loss: 1.6687
Epoch: 001/015 | Batch: 400/1000 | Loss: 1.5434
Epoch: 001/015 | Batch: 500/1000 | Loss: 1.2419
Epoch: 001/015 | Batch: 600/1000 | Loss: 1.3140
Epoch: 001/015 | Batch: 700/1000 | Loss: 1.6638
Epoch: 001/015 | Batch: 800/1000 | Loss: 1.1280
Epoch: 001/015 | Batch: 900/1000 | Loss: 1.4253
Epoch: 001/015 | Batch: 1000/1000 | Loss: 1.4685
Epoch: 001 Learning Rate 0.03000000 -> 0.03000000
Train Accuracy: 56.016
Validation Accuracy: 56.160
Test Accuracy: 56.000
Time elapsed so far: 0.78 min
Epoch: 002/015 | Batch: 100/1000 | Loss: 1.5867
Epoch: 002/015 | Batch: 200/1000 | Loss: 1.2620
Epoch: 002/015 | Batch: 300/1000 | Loss: 1.2938
Epoch: 002/015 | Batch: 400/1000 | Loss: 1.1416
Epoch: 002/015 | Batch: 500/1000 | Loss: 0.8978
Epoch: 002/015 | Batch: 600/1000 | Loss: 1.0315
Epoch: 002/015 | Batch: 700/1000 | Loss: 1.3266
Epoch: 002/015 | Batch: 800/1000 | Loss: 0.8658
Epoch: 002/015 | Batch: 900/1000 | Loss: 1.4317
Epoch: 002/015 | Batch: 1000/1000 | Loss: 1.2024
Epoch: 002 Learning Rate 0.03000000 -> 0.01500000
Train Accuracy: 65.932
Validation Accuracy: 66.360
Time elapsed so far: 1.51 min
Epoch: 003/015 | Batch: 100/1000 | Loss: 1.1343
Epoch: 003/015 | Batch: 200/1000 | Loss: 1.0333
Epoch: 003/015 | Batch: 300/1000 | Loss: 0.9750
Epoch: 003/015 | Batch: 400/1000 | Loss: 0.7618
Epoch: 003/015 | Batch: 500/1000 | Loss: 0.5869
Epoch: 003/015 | Batch: 600/1000 | Loss: 0.6248
Epoch: 003/015 | Batch: 700/1000 | Loss: 0.8990
Epoch: 003/015 | Batch: 800/1000 | Loss: 0.8073
Epoch: 003/015 | Batch: 900/1000 | Loss: 1.1276
Epoch: 003/015 | Batch: 1000/1000 | Loss: 1.0206
Epoch: 003 Learning Rate 0.01500000 -> 0.01500000
Train Accuracy: 73.822
Validation Accuracy: 73.820
Time elapsed so far: 2.22 min
Epoch: 004/015 | Batch: 100/1000 | Loss: 1.0532
Epoch: 004/015 | Batch: 200/1000 | Loss: 0.9303
Epoch: 004/015 | Batch: 300/1000 | Loss: 0.7329
Epoch: 004/015 | Batch: 400/1000 | Loss: 0.7059
Epoch: 004/015 | Batch: 500/1000 | Loss: 0.4901
Epoch: 004/015 | Batch: 600/1000 | Loss: 0.6889
Epoch: 004/015 | Batch: 700/1000 | Loss: 1.0142
Epoch: 004/015 | Batch: 800/1000 | Loss: 0.7893
Epoch: 004/015 | Batch: 900/1000 | Loss: 0.7046
Epoch: 004/015 | Batch: 1000/1000 | Loss: 0.8285
Epoch: 004 Learning Rate 0.01500000 -> 0.00750000
Train Accuracy: 76.318
Validation Accuracy: 76.040

Time elapsed so far: 2.94 min
Epoch: 005/015 | Batch: 100/1000 | Loss: 1.0377
Epoch: 005/015 | Batch: 200/1000 | Loss: 0.8578
Epoch: 005/015 | Batch: 300/1000 | Loss: 0.6407
Epoch: 005/015 | Batch: 400/1000 | Loss: 0.6001
Epoch: 005/015 | Batch: 500/1000 | Loss: 0.5147
Epoch: 005/015 | Batch: 600/1000 | Loss: 0.5106
Epoch: 005/015 | Batch: 700/1000 | Loss: 0.7989
Epoch: 005/015 | Batch: 800/1000 | Loss: 0.6127
Epoch: 005/015 | Batch: 900/1000 | Loss: 0.7235
Epoch: 005/015 | Batch: 1000/1000 | Loss: 0.8392
Epoch: 005 Learning Rate 0.00750000 -> 0.00750000
Train Accuracy: 79.652
Validation Accuracy: 80.320
Time elapsed so far: 3.67 min
Epoch: 006/015 | Batch: 100/1000 | Loss: 0.9017
Epoch: 006/015 | Batch: 200/1000 | Loss: 0.8267
Epoch: 006/015 | Batch: 300/1000 | Loss: 0.6443
Epoch: 006/015 | Batch: 400/1000 | Loss: 0.5671
Epoch: 006/015 | Batch: 500/1000 | Loss: 0.4894
Epoch: 006/015 | Batch: 600/1000 | Loss: 0.3716
Epoch: 006/015 | Batch: 700/1000 | Loss: 0.7554
Epoch: 006/015 | Batch: 800/1000 | Loss: 0.6736
Epoch: 006/015 | Batch: 900/1000 | Loss: 0.6263
Epoch: 006/015 | Batch: 1000/1000 | Loss: 0.6891
Epoch: 006 Learning Rate 0.00750000 -> 0.00375000
Train Accuracy: 80.872
Validation Accuracy: 81.780
Test Accuracy: 77.730
Time elapsed so far: 4.44 min
Epoch: 007/015 | Batch: 100/1000 | Loss: 0.7641
Epoch: 007/015 | Batch: 200/1000 | Loss: 0.7283
Epoch: 007/015 | Batch: 300/1000 | Loss: 0.5325
Epoch: 007/015 | Batch: 400/1000 | Loss: 0.5867
Epoch: 007/015 | Batch: 500/1000 | Loss: 0.4014
Epoch: 007/015 | Batch: 600/1000 | Loss: 0.4202
Epoch: 007/015 | Batch: 700/1000 | Loss: 0.7304
Epoch: 007/015 | Batch: 800/1000 | Loss: 0.5312
Epoch: 007/015 | Batch: 900/1000 | Loss: 0.5786
Epoch: 007/015 | Batch: 1000/1000 | Loss: 0.6558
Epoch: 007 Learning Rate 0.00375000 -> 0.00375000
Train Accuracy: 83.194
Validation Accuracy: 83.860
Time elapsed so far: 5.17 min
Epoch: 008/015 | Batch: 100/1000 | Loss: 0.8473
Epoch: 008/015 | Batch: 200/1000 | Loss: 0.5557
Epoch: 008/015 | Batch: 300/1000 | Loss: 0.3619
Epoch: 008/015 | Batch: 400/1000 | Loss: 0.6176
Epoch: 008/015 | Batch: 500/1000 | Loss: 0.3344
Epoch: 008/015 | Batch: 600/1000 | Loss: 0.3872
Epoch: 008/015 | Batch: 700/1000 | Loss: 0.6459
Epoch: 008/015 | Batch: 800/1000 | Loss: 0.4836
Epoch: 008/015 | Batch: 900/1000 | Loss: 0.4790
Epoch: 008/015 | Batch: 1000/1000 | Loss: 0.7204
Epoch: 008 Learning Rate 0.00375000 -> 0.00187500
Train Accuracy: 83.798

Validation Accuracy: 84.960
Time elapsed so far: 5.88 min
Epoch: 009/015 | Batch: 100/1000 | Loss: 0.6746
Epoch: 009/015 | Batch: 200/1000 | Loss: 0.5347
Epoch: 009/015 | Batch: 300/1000 | Loss: 0.4934
Epoch: 009/015 | Batch: 400/1000 | Loss: 0.4639
Epoch: 009/015 | Batch: 500/1000 | Loss: 0.3056
Epoch: 009/015 | Batch: 600/1000 | Loss: 0.3671
Epoch: 009/015 | Batch: 700/1000 | Loss: 0.5834
Epoch: 009/015 | Batch: 800/1000 | Loss: 0.4388
Epoch: 009/015 | Batch: 900/1000 | Loss: 0.5162
Epoch: 009/015 | Batch: 1000/1000 | Loss: 0.5077
Epoch: 009 Learning Rate 0.00187500 -> 0.00187500
Train Accuracy: 85.042
Validation Accuracy: 86.300
Time elapsed so far: 6.58 min
Epoch: 010/015 | Batch: 100/1000 | Loss: 0.6461
Epoch: 010/015 | Batch: 200/1000 | Loss: 0.5911
Epoch: 010/015 | Batch: 300/1000 | Loss: 0.4761
Epoch: 010/015 | Batch: 400/1000 | Loss: 0.4635
Epoch: 010/015 | Batch: 500/1000 | Loss: 0.3038
Epoch: 010/015 | Batch: 600/1000 | Loss: 0.2460
Epoch: 010/015 | Batch: 700/1000 | Loss: 0.5415
Epoch: 010/015 | Batch: 800/1000 | Loss: 0.5303
Epoch: 010/015 | Batch: 900/1000 | Loss: 0.5549
Epoch: 010/015 | Batch: 1000/1000 | Loss: 0.4717
Epoch: 010 Learning Rate 0.00187500 -> 0.00093750
Train Accuracy: 85.350
Validation Accuracy: 85.940
Time elapsed so far: 7.31 min
Epoch: 011/015 | Batch: 100/1000 | Loss: 0.6829
Epoch: 011/015 | Batch: 200/1000 | Loss: 0.5794
Epoch: 011/015 | Batch: 300/1000 | Loss: 0.4116
Epoch: 011/015 | Batch: 400/1000 | Loss: 0.5402
Epoch: 011/015 | Batch: 500/1000 | Loss: 0.2365
Epoch: 011/015 | Batch: 600/1000 | Loss: 0.2585
Epoch: 011/015 | Batch: 700/1000 | Loss: 0.3914
Epoch: 011/015 | Batch: 800/1000 | Loss: 0.7106
Epoch: 011/015 | Batch: 900/1000 | Loss: 0.5158
Epoch: 011/015 | Batch: 1000/1000 | Loss: 0.5201
Epoch: 011 Learning Rate 0.00093750 -> 0.00093750
Train Accuracy: 86.306
Validation Accuracy: 86.720
Test Accuracy: 81.060
Time elapsed so far: 8.08 min
Epoch: 012/015 | Batch: 100/1000 | Loss: 0.8055
Epoch: 012/015 | Batch: 200/1000 | Loss: 0.5166
Epoch: 012/015 | Batch: 300/1000 | Loss: 0.4148
Epoch: 012/015 | Batch: 400/1000 | Loss: 0.4014
Epoch: 012/015 | Batch: 500/1000 | Loss: 0.2833
Epoch: 012/015 | Batch: 600/1000 | Loss: 0.2760
Epoch: 012/015 | Batch: 700/1000 | Loss: 0.5072
Epoch: 012/015 | Batch: 800/1000 | Loss: 0.4418
Epoch: 012/015 | Batch: 900/1000 | Loss: 0.4441
Epoch: 012/015 | Batch: 1000/1000 | Loss: 0.3717
Epoch: 012 Learning Rate 0.00093750 -> 0.00046875

Train Accuracy: 86.440
 Validation Accuracy: 87.560
 Time elapsed so far: 8.78 min
 Epoch: 013/015 | Batch: 100/1000 | Loss: 0.7890
 Epoch: 013/015 | Batch: 200/1000 | Loss: 0.5606
 Epoch: 013/015 | Batch: 300/1000 | Loss: 0.3905
 Epoch: 013/015 | Batch: 400/1000 | Loss: 0.3948
 Epoch: 013/015 | Batch: 500/1000 | Loss: 0.3171
 Epoch: 013/015 | Batch: 600/1000 | Loss: 0.2461
 Epoch: 013/015 | Batch: 700/1000 | Loss: 0.4524
 Epoch: 013/015 | Batch: 800/1000 | Loss: 0.3789
 Epoch: 013/015 | Batch: 900/1000 | Loss: 0.3493
 Epoch: 013/015 | Batch: 1000/1000 | Loss: 0.3727
 Epoch: 013 Learning Rate 0.00046875 -> 0.00046875
 Train Accuracy: 86.846
 Validation Accuracy: 86.980
 Time elapsed so far: 9.48 min
 Epoch: 014/015 | Batch: 100/1000 | Loss: 0.7272
 Epoch: 014/015 | Batch: 200/1000 | Loss: 0.4903
 Epoch: 014/015 | Batch: 300/1000 | Loss: 0.4136
 Epoch: 014/015 | Batch: 400/1000 | Loss: 0.4935
 Epoch: 014/015 | Batch: 500/1000 | Loss: 0.1816
 Epoch: 014/015 | Batch: 600/1000 | Loss: 0.2322
 Epoch: 014/015 | Batch: 700/1000 | Loss: 0.4211
 Epoch: 014/015 | Batch: 800/1000 | Loss: 0.5411
 Epoch: 014/015 | Batch: 900/1000 | Loss: 0.4382
 Epoch: 014/015 | Batch: 1000/1000 | Loss: 0.4898
 Epoch: 014 Learning Rate 0.00046875 -> 0.00023437
 Train Accuracy: 86.966
 Validation Accuracy: 88.160
 Time elapsed so far: 10.20 min
 Epoch: 015/015 | Batch: 100/1000 | Loss: 0.6864
 Epoch: 015/015 | Batch: 200/1000 | Loss: 0.4992
 Epoch: 015/015 | Batch: 300/1000 | Loss: 0.3462
 Epoch: 015/015 | Batch: 400/1000 | Loss: 0.4666
 Epoch: 015/015 | Batch: 500/1000 | Loss: 0.2510
 Epoch: 015/015 | Batch: 600/1000 | Loss: 0.3214
 Epoch: 015/015 | Batch: 700/1000 | Loss: 0.5002
 Epoch: 015/015 | Batch: 800/1000 | Loss: 0.4640
 Epoch: 015/015 | Batch: 900/1000 | Loss: 0.4818
 Epoch: 015/015 | Batch: 1000/1000 | Loss: 0.4726
 Epoch: 015 Learning Rate 0.00023437 -> 0.00023437
 Train Accuracy: 86.996
 Validation Accuracy: 87.420
 Time elapsed so far: 10.91 min
 Total Train Time: 10.91 min

```

In [ ]: model = model.eval()
        ts_acc = accuracy(model, testloader, device)
        print(f"Test Accuracy: {ts_acc:0.3f}")
  
```

Test Accuracy: 81.000

```

In [ ]: # torch.save(model.state_dict(), 'model_state_dict.pth')
  
```

Evaluation:

```
In [ ]: model.eval()

#Tensors for ypred and ytrue
ypred = torch.tensor([], dtype=torch.long, device=device)
ytrue = torch.tensor([], dtype=torch.long, device=device)

with torch.no_grad():
    for imgs, labels in testloader:
        imgs = imgs.to(device)
        labels = labels.to(device)

        logits = model(imgs)
        _, label_pred = torch.max(logits, axis=1)

        ypred = torch.cat((ypred, label_pred))
        ytrue = torch.cat((ytrue, labels))
ypred = ypred.cpu().numpy()
ytrue = ytrue.cpu().numpy()
```

Visualizations:

```
In [ ]: cm = confusion_matrix(ytrue,ypred,labels=[*range(10)])
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.title("Confusion Matrix for model")
plt.show()

#Line plot for loss
plt.plot([*range(1,num_epochs+1)],[x.cpu().detach().numpy() for x in train_loss])
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.xticks([*range(1,num_epochs+1)])
plt.title("Loss per Epoch")
plt.show()
```

